

VectorEditGym: Evaluating Specification-Faithful SVG Repair Under Preservation Constraints

Yug Aditi Gupta

Theta Labs

yug@thetalab.tech

Abstract

Instruction-based vector editing requires two capabilities: making a requested change and leaving everything else alone. The second is easy to miss when an output is judged only as a raster image. We introduce **VectorEditGym**, a compact, difficult benchmark of 40 SVG repair tasks. Each task pairs a corrupted SVG program with a human visual instruction, a hidden target program, 5.05 annotated repairs on average, and an average of 60.55 protected objects. Instructions describe visible defects without exposing element identifiers, coordinates, color codes, or path data. We define a deterministic binary reward that is one only when the output parses, implements every annotated repair, and is canonically equivalent to the target; edit completion and Unintended Change Rate (UCR) diagnose partial outcomes. We evaluate 30 model endpoints (25 listed as open-weight and 5 inexpensive controls) over 1200 requests. The strongest endpoint reaches only 2.5% binary success, despite 40.7% requested-edit completion, showing that apparent repair progress and specification-faithful editing remain substantially different. All prompts, outputs, scoring code, costs, and per-task reports are released.

1 Introduction

An instruction such as “the amber signal has gone dark; fix it and leave the station alone” defines both a positive requirement and a large set of negative requirements. The signal must change, while the train, clock, cables, platform, and every unrelated program detail must not. A model can produce a plausible raster while renaming elements, rewriting path data, shifting an unrelated object, or deleting metadata. These are not cosmetic differences for an SVG: they alter an executable, structured artifact (World Wide Web Consortium, 2018).

Most generation metrics emphasize final appearance. That is appropriate for open-ended synthesis,

but insufficient for local editing. A repair benchmark must answer three separate questions: Did the requested objects change? Did unmentioned objects remain stable? Is the resulting program valid and usable? VectorEditGym makes all three observable through hidden target SVGs and explicit protected-object sets.

The benchmark is deliberately small and dense rather than broad and shallow. Its 40 tasks contain 202 annotated repairs across station, harbor, campsite, market, and larger scenic illustrations. The public instruction never names a DOM identifier or serialized numeric/path value. Solvers must interpret the visible scene and edit the source accordingly. Figure 1 shows representative inputs and targets.

Our contributions are:

1. a frozen corpus of dense, human-described SVG repairs with explicit intended edits and broad preservation surfaces;
2. a transparent executable evaluator with binary reward, edit completion, validity, preservation, and UCR;
3. a reproducible 30-endpoint evaluation with response provenance, token usage, cost, failures, and model-produced SVGs; and
4. an empirical account of the gap between completing visible repairs and preserving the full vector program.

2 Related Work

Vector generation. IconShop generates vector icons from text (Wu et al., 2023); StarVector generates SVG code from images and text (Rodriguez et al., 2025); and LLM4SVG studies complex vector understanding and generation with language

Corrupted input




Hidden target




The amber signal has gone dark, the train door is sitting too far to the right, and the overhead cable has an awkward sag. Put those three details back in place and leave the rest of the station alone.

Corrupted input




Hidden target




In the deep-sea scene, remove the stray colored marker and move the displaced glowing particle back toward the nearby cluster. The large jellyfish should have fine tentacles and a soft pink lower rim, while the second shaft of light should be clearly visible. Leave the other creatures alone.

Corrupted input



Hidden target



Remove the colored marker from the Venetian sunset. Center the large sun glow behind the canal, restore the dark red body of the gondola, slim the main arch of the bridge, and bring back the warm glow in the affected window. Keep the other buildings and reflections intact.

Figure 1: Representative corrupted inputs and hidden targets. Public prompts describe visible mistakes in ordinary language; source IDs and serialized numeric/path values remain evaluator-only annotations.

models (Xing et al., 2025). These systems establish SVG as a useful meeting point between visual generation and program synthesis. VectorEditGym focuses on local repair rather than open-ended generation.

Instruction-based SVG editing. SVGEditBench introduced quantitative evaluation for LLM-based SVG editing (Nishina and Matsui, 2024), and SVGEditBench V2 expanded instruction-based editing evaluation (Nishina and Matsui, 2025). VectorEdits provides a larger dataset and benchmark for vector-graphics editing (Kuchař et al., 2025). Our scope is complementary: we emphasize dense multi-repair scenes, strict program equivalence after representation normalization, per-object preservation, and publication of every model output. We do not claim that a single structural target covers all visually

acceptable edits; instead, we use it to measure specification-faithful repair under an explicitly strict contract.

3 Benchmark

3.1 Task Format

Each task is a tuple

$$\tau = (x, S_0, S^*, E, U),$$

where x is a human repair instruction, S_0 is the corrupted SVG, S^* is the hidden target, E is a set of expected object-attribute repairs, and U is the set of object IDs that must be preserved. The model receives only (x, S_0) . Target source, expected diffs, target-part lists, and preserve lists are excluded from the prompt and recorded as hidden evaluator inputs.

Instructions were rewritten after task freezing. They describe visible objects and relations—for

Statistic	Value
Tasks	40
Annotated repairs	202
Repairs per task (mean)	5.05
Protected objects per task (mean)	60.55
SVG elements per task (mean)	85.4
Input characters (total)	294,798

Table 1: Frozen corpus statistics. Element counts include non-rendering SVG nodes; protected-object counts refer to identified scene objects.




Figure 2: Repair operations and per-task edit density.

example, a moon that drifted down and left or a jellyfish whose tentacles became too heavy—without SVG IDs, numeric coordinates, hex colors, path commands, or attribute names. A corpus regression hash covers every non-instruction task field, ensuring this language-only rewrite did not change inputs, targets, annotations, or ordering.

3.2 Composition and Difficulty

The corpus contains 20 compact but busy scenes derived from four authored scene templates and 20 larger scenic illustrations. Ten tasks are marked hard and 30 very hard. Tasks require between three and eight repairs, including insertion or deletion, recoloring, displacement, path restoration, line-weight correction, and opacity restoration. All unrequested identified parts are protected.

4 Executable Evaluation

4.1 Canonical Program Comparison

Let $C(S)$ parse an SVG and canonicalize representation-only variation. We discard namespace spelling differences and attribute order, normalize numeric spellings and separators in paths, points, transforms, and view boxes, and map common equivalent color spellings (for example, white, #fff, and rgb(255,255,255)) to one form. Element order, nesting, IDs, geometry, paint, text, and other program semantics remain strict. Malformed XML has no canonical form.




Figure 3: SVG size and preservation surface. Marker area scales with the number of requested repairs.

The primary reward is binary:

$$R(S) = \mathbb{1}[C(S) = C(S^*)].$$

Thus a solver receives one point only if every requested repair and every preservation constraint is satisfied. There is no learned judge, reference selection heuristic or synthetic reference agent in the public protocol.

4.2 Diagnostic Metrics

For expected checks E , edit completion is

$$\text{Edit}(S) = \frac{1}{|E|} \sum_{e \in E} \mathbb{1}[e(S) = e(S^*)].$$

For protected objects U , preservation and unintended-change rate are

$$\text{Pres}(S) = \frac{1}{|U|} \sum_{u \in U} \mathbb{1}[C(u_S) = C(u_{S^*})],$$

$$\text{UCR}(S) = 1 - \text{Pres}(S).$$

We additionally report parse validity, normalized source exact match, canonical structural match, provider errors, scheduled elapsed time, tokens, and cost. A provider error or malformed output receives reward zero. Diagnostic metrics never replace the binary primary reward.

5 Model Study

5.1 Protocol

We evaluate 25 endpoints designated as open-weight in the study manifest and 5 inexpensive closed controls through the OpenRouter chat-completions API (OpenRouter, 2026). This designation records the endpoint grouping used for analysis and is not an independent license audit. The exact endpoint IDs appear in Appendix A.

Every model receives the same system instruction, human repair request, and corrupted SVG source. We record one scored outcome per task, omit temperature rather than imposing a provider-incompatible value, and adapt the output ceiling to input length (minimum 4,096; maximum 12,000 tokens). The runner retries rate limits and transient server failures against the same endpoint. It never falls back to another model; the client permits two transport retries inside each harness attempt. We record requested and resolved model IDs, response ID, finish reason, prompt, completion and reasoning tokens when supplied, scheduled wall-clock elapsed time, and provider-reported cost. Elapsed time begins when a model-task job is scheduled and therefore includes local concurrency-queue wait; we expose it for run auditing, not as endpoint latency. A reservation-based budget guard caps the run at \$25; the completed study cost \$7.62.

5.2 Uncertainty

We compute 95% confidence intervals by resampling the 40 tasks with replacement 10,000 times independently for each model. These intervals quantify task-set uncertainty only. They do not estimate sampling variance because each model-task pair has one scored outcome and no repeated decoding. Endpoint tables are ordered by binary reward, then edit completion, lower UCR, and model name.

6 Results

Binary success remains rare. The strongest endpoint, Gemma 4 31B IT, obtains 2.5% binary reward. Its edit completion is substantially higher at 40.7%, while its UCR is 10.8% and its parse validity is 90.0%. The difference is the central benchmark result: models frequently perform some requested visual work but fail the full repair-and-preserve contract. Figure 4 reports task-bootstrap intervals for every endpoint.

Repair and preservation are separate axes. Figure 5 plots requested-edit completion against UCR. A model can move right by correcting visible defects while also moving upward by rewriting protected scene objects. Reporting completion alone would hide this behavior; reporting only full success would hide useful partial capability.

Operation families expose different bottlenecks. Figure 6 decomposes the expected

checks. Simple color or add/remove corrections need not predict path-shape or position restoration, especially because the prompt gives visual descriptions instead of literal DOM values. This decomposition also distinguishes an endpoint that attempted an edit but selected the wrong source element from one that returned an invalid or unchanged document.

Cost is not a sufficient capability proxy. The run spans inexpensive small endpoints and larger reasoning or coding models. Figure 7 shows edit completion against total 40-task cost. It is a descriptive endpoint comparison rather than a hardware-efficiency claim: OpenRouter pricing and routing can change, and the study uses one scored outcome per task.

7 Error Analysis

Literal partial repair. Models often identify an obvious semantic defect, such as a signal with the wrong color, while leaving a displaced object or malformed curve unresolved. These outputs earn partial edit-completion credit but binary reward zero.

Global rewrites. Some responses regenerate or reformat substantial portions of the SVG. Canonical normalization accepts harmless formatting, but changed paths, reordered scene objects, renamed IDs, or missing definitions count as program changes. Protected object checks identify which named scene components drifted; document-level checks catch anonymous definitions and root attributes.

Ambiguous visual recovery. Human prompts intentionally omit exact coordinates and path commands. A model may make a visually reasonable adjustment that differs from the frozen target. The strict reward records that as failure. This is useful for evaluating exact repair of structured assets but should not be interpreted as a complete perceptual quality judgment.

Provider and truncation failures. Only 56.9% of scored outcomes contain parseable SVG. The harness preserves raw malformed output, API errors, and finish reasons in the released records. Transient failures are retried only against the requested endpoint; persistent failures remain zero-scored observations. The aggregate provider error rate is 3.4%; parse failure is measured separately.

Model	Reward $\uparrow$	Edit $\uparrow$	UCR $\downarrow$	Valid $\uparrow$	Cost
Gemma 4 31B IT	2.5	40.7	10.8	90.0	\$0.084
Qwen3 Coder	0.0	38.9	1.4	100.0	\$0.295
HY 3	0.0	38.8	1.8	100.0	\$0.092
GPT-OSS 120B	0.0	38.4	6.0	95.0	\$0.057
Seed 2.0 Mini	0.0	36.6	9.4	92.5	\$0.174
Devstral 2512	0.0	36.1	0.7	100.0	\$0.358
Gemini 3.1 Flash Lite	0.0	35.1	2.5	100.0	\$0.269
Llama 4 Maverick	0.0	34.6	7.4	95.0	\$0.143
Mistral Small 2603	0.0	31.4	6.5	100.0	\$0.119
GPT-OSS 20B	0.0	29.9	28.7	72.5	\$0.030

Table 2: Top ten endpoints under binary reward. Reward, edit completion, and UCR and validity are percentages; cost is the provider-reported cost for all 40 tasks. Full results are in Appendix A.

8 Limitations and Artifact Risks

VectorEditGym has only 40 tasks and should be treated as a focused stress test, not a comprehensive measure of visual editing. Public targets make future contamination possible. The single-target structural reward can reject visually or semantically acceptable alternatives, and our basic color canonicalization is not a complete CSS equivalence engine. Conversely, canonical tree equality does not test downstream animation, scripting, accessibility, or editor-specific behavior.

The evaluation uses one scored outcome per model-task pair and therefore cannot characterize sampling stability. OpenRouter can change prices, routing, model versions, and availability; requested and resolved IDs are published to make this visible. The “open-weight” grouping is manifest metadata, not a legal determination.

Twenty scenic source files in the frozen corpus lack recoverable per-file provenance metadata. We do not claim original authorship of those illustrations. They are retained to preserve the evaluated task set, marked as provenance pending in the artifact, and should be treated as research-only until provenance is resolved. This limits redistribution confidence and is a material artifact risk.

Finally, the benchmark evaluates text-and-source editing rather than multimodal screen interaction. Models see SVG code and a visual description, but not a raster preview or interactive editor. Performance therefore combines scene reasoning, source localization, and code generation under one protocol.

9 Conclusion

VectorEditGym isolates a practical requirement for editing agents: requested change is not enough; preservation is part of the specification. Human

visual instructions make the task less like copying a patch, while executable SVG targets make every decision auditable. Across 30 endpoints, partial repair is much more common than exact repair-and-preserve success. The released outputs and diff reports make that gap inspectable at the object and source level and provide a compact target for improving structured visual editors.

Reproducibility Statement

The repository includes the frozen corpus hash, model manifest, prompt template, budgeted resumable runner, canonical evaluator, unit tests, raw-result schema, analysis script, generated tables, figures, paper source, Harbor tasks, and a web viewer for every model output. API credentials and ignored raw working directories are not included.




Figure 4: Binary task success for all endpoints. Orange bars are inexpensive controls; teal bars are endpoints listed as open-weight in the manifest.




Figure 5: Requested-edit completion versus unintended-change rate. Marker area scales with binary reward. Lower right is preferable.




Figure 6: Expected-edit check success by operation family. Rows follow overall binary-reward rank.




Figure 7: Edit completion versus provider-reported run cost.

A Full Endpoint Results

Values are percentages. Errors include persistent provider failures and missing usable SVG output. Endpoint IDs, resolved IDs, response IDs, and per-task details are available in the machine-readable release.

Model	Group	Reward	Edit	UCR	Valid	Errors
Gemma 4 31B IT	open-weight	2.5	40.7	10.8	90.0	0.0
Qwen3 Coder	open-weight	0.0	38.9	1.4	100.0	0.0
HY 3	open-weight	0.0	38.8	1.8	100.0	0.0
GPT-OSS 120B	open-weight	0.0	38.4	6.0	95.0	0.0
Seed 2.0 Mini	cheap-control	0.0	36.6	9.4	92.5	0.0
Devstral 2512	open-weight	0.0	36.1	0.7	100.0	0.0
Gemini 3.1 Flash Lite	cheap-control	0.0	35.1	2.5	100.0	0.0
Llama 4 Maverick	open-weight	0.0	34.6	7.4	95.0	0.0
Mistral Small 2603	open-weight	0.0	31.4	6.5	100.0	0.0
GPT-OSS 20B	open-weight	0.0	29.9	28.7	72.5	0.0
Qwen3.5 397B A17B	open-weight	0.0	28.7	35.9	65.0	0.0
GPT-5.4 Nano	cheap-control	0.0	26.7	2.2	100.0	0.0
Phi-4	open-weight	0.0	24.9	3.7	100.0	0.0
MiniMax M3	open-weight	0.0	24.5	55.4	45.0	0.0
DeepSeek V4 Pro	open-weight	0.0	23.2	53.1	47.5	0.0
Inkling	open-weight	0.0	20.6	55.4	45.0	0.0
KAT Coder Air V2.5	cheap-control	0.0	19.6	57.6	42.5	0.0
Granite 4.1 8B	open-weight	0.0	18.4	20.6	95.0	0.0
Laguna M.1	open-weight	0.0	16.8	58.3	42.5	0.0
Ling 2.6 Flash	cheap-control	0.0	15.7	10.1	95.0	0.0
Kimi K2.6	open-weight	0.0	13.2	75.1	25.0	2.5
GLM 4.7	open-weight	0.0	10.7	70.3	30.0	0.0
Nemotron 3 Ultra	open-weight	0.0	3.8	92.5	7.5	0.0
Qwen3.6 35B A3B	open-weight	0.0	3.8	92.6	7.5	0.0
MiMo V2.5 Pro	open-weight	0.0	3.1	90.1	10.0	0.0
Nemotron 3 Nano	open-weight	0.0	1.2	97.5	2.5	0.0
Trinity Large Thinking	open-weight	0.0	1.2	97.5	2.5	0.0
OLMo 3 32B Think	open-weight	0.0	0.0	100.0	0.0	100.0
Reka Flash 3	open-weight	0.0	0.0	100.0	0.0	0.0
Step 3.7 Flash	open-weight	0.0	0.0	100.0	0.0	0.0

References

- Josef Kuchař, Marek Kadlčík, Michal Spiegel, and Michal Štefánik. 2025. [VectorEdits: A dataset and benchmark for instruction-based editing of vector graphics](#). *Preprint*, arXiv:2506.15903.
- Kunato Nishina and Yusuke Matsui. 2024. [SVGEdit-Bench: A benchmark dataset for quantitative assessment of LLM’s SVG editing capabilities](#). *Preprint*, arXiv:2404.13710.
- Kunato Nishina and Yusuke Matsui. 2025. [SVGEdit-Bench V2: A benchmark for instruction-based SVG editing](#). *Preprint*, arXiv:2502.19453.
- OpenRouter. 2026. Api reference and model catalog. <https://openrouter.ai/docs/api/reference/overview>. Accessed 2026-07-20.
- Juan A. Rodriguez, Abhay Puri, Shubham Agarwal, Issam H. Laradji, Pau Rodriguez, Sai Rajeswar, David Vazquez, Christopher Pal, and Marco Pedersoli. 2025. [StarVector: Generating scalable vector graphics code from images and text](#). *Preprint*, arXiv:2312.11556.
- World Wide Web Consortium. 2018. [Scalable vector graphics \(SVG\) 2](#). Technical report, W3C.
- Ronghuan Wu, Wanchao Su, Kede Ma, and Jing Liao. 2023. [IconShop: Text-guided vector icon synthesis with autoregressive transformers](#). *Preprint*, arXiv:2304.14400.
- Ximing Xing, Juncheng Hu, Guotao Liang, Jing Zhang, Dong Xu, and Qian Yu. 2025. [Empowering LLMs to understand and generate complex vector graphics](#). *Preprint*, arXiv:2412.11102.